



Sardar Patel Institute of Technology, Mumbai  
Department of Electronics and Telecommunication Engineering  
B.E. Sem-VII

### ISE-2 ML Minor

**Name: Sruthi Shivaramakrishnan    Batch:A    UID:2019110059    Branch: ETRX**

**Objective: Dealing with imbalanced dataset (Given on Kaggle data repository)**

The dataset contains transactions made by credit cards in September 2013 by European cardholders. This dataset presents transactions that occurred in two days, where you may observe 492 frauds out of 284,807 transactions. The dataset is highly unbalanced, the positive class (frauds) account for 0.172% of all transactions.

- Understand the imbalanced dataset and will perform various approaches like undersampling/ oversampling, choosing the right metrics of ROC- AUC to deal with the imbalanced dataset.
- After trying different approaches and training different models (LR, KNN, SVM, Random Forest, XGboost, Naive Bayes) you will compare their results and decide the one which fits best for your application.

**Platform : Google Colab**

**Data Set : Credit Card Fraud Detection**

**Code and output:**

```
# Importing libraries
```

```
from google.colab import drive  
drive.mount('/content/drive')
```

```
import pandas as pd
```

```
# Reading the dataset
```

```
df=pd.read_csv('/content/drive/MyDrive/creditcard.csv')
```

```
df.head()
```

|   | Time | V1        | V2        | V3       | V4        | V5        | V6        | V7        | V8        | V9        | ... | V21       | V22       | V23       | V24       |
|---|------|-----------|-----------|----------|-----------|-----------|-----------|-----------|-----------|-----------|-----|-----------|-----------|-----------|-----------|
| 0 | 0.0  | -1.359807 | -0.072781 | 2.536347 | 1.378155  | -0.338321 | 0.462388  | 0.239599  | 0.098698  | 0.363787  | ... | -0.018307 | 0.277838  | -0.110474 | 0.066928  |
| 1 | 0.0  | 1.191857  | 0.266151  | 0.166480 | 0.448154  | 0.060018  | -0.082361 | -0.078803 | 0.085102  | -0.255425 | ... | -0.225775 | -0.638672 | 0.101288  | -0.339846 |
| 2 | 1.0  | -1.358354 | -1.340163 | 1.773209 | 0.379780  | -0.503198 | 1.800499  | 0.791461  | 0.247676  | -1.514654 | ... | 0.247998  | 0.771679  | 0.909412  | -0.689281 |
| 3 | 1.0  | -0.966272 | -0.185226 | 1.792993 | -0.863291 | -0.010309 | 1.247203  | 0.237609  | 0.377436  | -1.387024 | ... | -0.108300 | 0.005274  | -0.190321 | -1.175575 |
| 4 | 2.0  | -1.158233 | 0.877737  | 1.548718 | 0.403034  | -0.407193 | 0.095921  | 0.592941  | -0.270533 | 0.817739  | ... | -0.009431 | 0.798278  | -0.137458 | 0.141267  |

5 rows × 31 columns

# Checking null values

df.isnull().sum()

```
Time      0
V1        0
V2        0
V3        0
V4        0
V5        0
V6        0
V7        0
V8        0
V9        0
V10       0
V11       0
V12       0
V13       0
V14       0
V15       0
V16       0
V17       0
V18       0
V19       0
V20       0
V21       0
V22       0
V23       0
V24       0
V25       0
V26       0
V27       0
V28       0
Amount    0
Class     0
dtype: int64
```

**Data has no values**

df.shape

```
(284807, 31)
```

df.describe()

|       | Time          | V1            | V2            | V3            | V4            | V5            | V6            | V7            | V8            |          |
|-------|---------------|---------------|---------------|---------------|---------------|---------------|---------------|---------------|---------------|----------|
| count | 284807.000000 | 2.848070e+05  | 2.848070e+05  | 2.848070e+05  | 2.848070e+05  | 2.848070e+05  | 2.848070e+05  | 2.848070e+05  | 2.848070e+05  | 2.84807  |
| mean  | 94813.859575  | 1.168375e-15  | 3.416908e-16  | -1.379537e-15 | 2.074095e-15  | 9.604066e-16  | 1.487313e-15  | -5.556467e-16 | 1.213481e-16  | -2.40633 |
| std   | 47488.145955  | 1.958696e+00  | 1.651309e+00  | 1.516255e+00  | 1.415869e+00  | 1.380247e+00  | 1.332271e+00  | 1.237094e+00  | 1.194353e+00  | 1.09863  |
| min   | 0.000000      | -5.640751e+01 | -7.271573e+01 | -4.832559e+01 | -5.683171e+00 | -1.137433e+02 | -2.616051e+01 | -4.355724e+01 | -7.321672e+01 | -1.34340 |
| 25%   | 54201.500000  | -9.203734e-01 | -5.985499e-01 | -8.903648e-01 | -8.486401e-01 | -6.915971e-01 | -7.682956e-01 | -5.540759e-01 | -2.086297e-01 | -6.43097 |
| 50%   | 84692.000000  | 1.810880e-02  | 6.548556e-02  | 1.798463e-01  | -1.984653e-02 | -5.433583e-02 | -2.741871e-01 | 4.010308e-02  | 2.235804e-02  | -5.14287 |
| 75%   | 139320.500000 | 1.315642e+00  | 8.037239e-01  | 1.027196e+00  | 7.433413e-01  | 6.119264e-01  | 3.985649e-01  | 5.704361e-01  | 3.273459e-01  | 5.97135  |
| max   | 172792.000000 | 2.454930e+00  | 2.205773e+01  | 9.382558e+00  | 1.687534e+01  | 3.480167e+01  | 7.330163e+01  | 1.205895e+02  | 2.000721e+01  | 1.55949  |

8 rows × 31 columns

df.info()

| #  | Column | Non-Null Count  | Dtype   |
|----|--------|-----------------|---------|
| 0  | Time   | 284807 non-null | float64 |
| 1  | V1     | 284807 non-null | float64 |
| 2  | V2     | 284807 non-null | float64 |
| 3  | V3     | 284807 non-null | float64 |
| 4  | V4     | 284807 non-null | float64 |
| 5  | V5     | 284807 non-null | float64 |
| 6  | V6     | 284807 non-null | float64 |
| 7  | V7     | 284807 non-null | float64 |
| 8  | V8     | 284807 non-null | float64 |
| 9  | V9     | 284807 non-null | float64 |
| 10 | V10    | 284807 non-null | float64 |
| 11 | V11    | 284807 non-null | float64 |
| 12 | V12    | 284807 non-null | float64 |
| 13 | V13    | 284807 non-null | float64 |
| 14 | V14    | 284807 non-null | float64 |
| 15 | V15    | 284807 non-null | float64 |
| 16 | V16    | 284807 non-null | float64 |
| 17 | V17    | 284807 non-null | float64 |
| 18 | V18    | 284807 non-null | float64 |
| 19 | V19    | 284807 non-null | float64 |
| 20 | V20    | 284807 non-null | float64 |
| 21 | V21    | 284807 non-null | float64 |
| 22 | V22    | 284807 non-null | float64 |
| 23 | V23    | 284807 non-null | float64 |
| 24 | V24    | 284807 non-null | float64 |
| 25 | V25    | 284807 non-null | float64 |
| 26 | V26    | 284807 non-null | float64 |
| 27 | V27    | 284807 non-null | float64 |
| 28 | V28    | 284807 non-null | float64 |
| 29 | Amount | 284807 non-null | float64 |
| 30 | Class  | 284807 non-null | int64   |

import seaborn as sns

correlation=df.corr()

import matplotlib.pyplot as plt

plt.figure(figsize=(12,10))

sns.heatmap(correlation,annot=True,cmap=plt.cm.Red)

[illegible]

```

V2      Amount  -0.531409
Amount  V2      -0.531409
Time     Time     1.000000
V15      V15      1.000000
V28      V28      1.000000
V1       V1       1.000000
V2       V2       1.000000
V3       V3       1.000000
V4       V4       1.000000
V5       V5       1.000000
V6       V6       1.000000
V7       V7       1.000000
V8       V8       1.000000
V9       V9       1.000000
V10      V10      1.000000
V11      V11      1.000000
V12      V12      1.000000
Amount   Amount   1.000000
V13      V13      1.000000
V16      V16      1.000000
V17      V17      1.000000
V18      V18      1.000000
V19      V19      1.000000
V20      V20      1.000000
V21      V21      1.000000
V22      V22      1.000000
V23      V23      1.000000
V24      V24      1.000000
V25      V25      1.000000
V26      V26      1.000000
V27      V27      1.000000
V14      V14      1.000000
Class    Class    1.000000
dtype: float64

```

df.columns

```

Index(['Time', 'V1', 'V2', 'V3', 'V4', 'V5', 'V6', 'V7', 'V8', 'V9', 'V10',
      'V11', 'V12', 'V13', 'V14', 'V15', 'V16', 'V17', 'V18', 'V19', 'V20',
      'V21', 'V22', 'V23', 'V24', 'V25', 'V26', 'V27', 'V28', 'Amount',
      'Class'],
      dtype='object')

```

```

x=df.iloc[:, 0:30].values
y=df['Class']

```

#Splitting into train and test set

```

from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test= train_test_split(x, y, test_size= 0.25, random_state=0)

```


```

import seaborn as sns
from sklearn.datasets import make_classification

x, y = make_classification(n_samples=5000, n_features=2, n_informative=2,
                          n_redundant=0, n_repeated=0, n_classes=3,
                          n_clusters_per_class=1,
                          weights=[0.01, 0.05],
                          class_sep=0.8, random_state=0)

import matplotlib.pyplot as plt
colors = ['#ef8a62' if v == 0 else '#f7f7f7' if v == 1 else '#67a9cf' for v in y]
kwargs_params = {'linewidth': 1, 'edgecolor': 'black'}
fig = plt.figure(figsize=(12,6))
plt.scatter(x[:, 0], x[:, 1], c=colors, **kwargs_params)
sns.despine()

```



**The above plot shows the bias and imbalance of the data points. The fraud class account for only a small portion of the data**

```

#Applying the XCG algorithm on the given data

# import library
from xgboost import XGBClassifier

xgb_model = XGBClassifier().fit(x_train, y_train)

# predict
xgb_y_predict = xgb_model.predict(x_test)

```

```

from sklearn.metrics import accuracy_score, precision_score, recall_score

xgb_score = accuracy_score(xgb_y_predict, y_test)

print('Accuracy score is:', xgb_score)

```

```
Accuracy score is: 0.9904
```

**The above code shows high accuracy. The reason for this is high imbalance in the data. Though the accuracy is high, because of the imbalance it might not be able to perform well in world cases of fraud detection.**

```

class_count_0, class_count_1 = df['Class'].value_counts()

# Separate class
class_0 = df[df['Class'] == 0]
class_1 = df[df['Class'] == 1] # print the shape of the class
print('class 0:', class_0.shape)
print('class 1:', class_1.shape)

```

```

class 0: (284315, 33)
class 1: (492, 33)

```

**The above snippet shows the number of data points in each class.**

#Isolation forest to detect anomalies

```

import numpy as np
from sklearn.ensemble import IsolationForest

random_state = np.random.RandomState(42)
model=IsolationForest(n_estimators=100,max_samples='auto',contamination=float(0.2),random_
state=random_state)

model.fit(x)
pred = model.predict(x)


pred_scores = -1*model.score_samples(x)

print(model.get_params())

plt.scatter(x[:, 0], x[:, 1], c=pred_scores, cmap='RdBu')

```

```
plt.colorbar(label='Simplified Anomaly Score')
plt.show()
```



**The above snippet shows other anomalies in the data. The large anomalies may be due to the imbalance.**

```
# Over sampling using SMOTE
```

```
# import library
from imblearn.over_sampling import SMOTE
from collections import Counter
```

```
smote = SMOTE()
```

```
# fit predictor and target variable
x_smote, y_smote = smote.fit_resample(x, y)
```

```
print('Original dataset shape', Counter(y))
print('Resample dataset shape', Counter(y_smote))
```

```
Original dataset shape Counter({0: 284315, 1: 492})
Resample dataset shape Counter({0: 284315, 1: 284298})
```

**To increase the number of fraud cases we are using SMOTE to increase the number of samples of the imbalanced data.**

```
# Over sampling using Adasyn
```

```
from imblearn.over_sampling import ADASYN
```

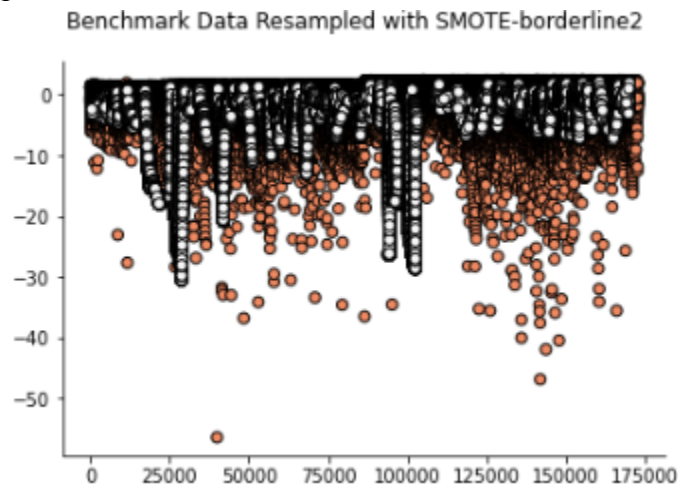
```
ada = ADASYN(random_state=42)
```



```

X_res, y_res = ada.fit_resample(x, y)
kwarg_params = {'linewidth': 1, 'edgecolor': 'black'}
colors = ['#ef8a62' if v == 0 else '#f7f7f7' if v == 1 else '#67a9cf' for v in y_res]
plt.scatter(X_res[:, 0], X_res[:, 1], c=colors, **kwargs)
sns.despine()
plt.suptitle("Benchmark Data Resampled with SMOTE-borderline2")
pass

```



```

print('Original dataset shape', Counter(y))
print('Resample dataset shape', Counter(y_res))

```

```

Original dataset shape Counter({0: 284315, 1: 492})
Resample dataset shape Counter({0: 284315, 1: 284298})

```

# Under sampling

# Random Under Sampling for Adasyn data

```
from imblearn.under_sampling import RandomUnderSampler
```

```

rus = RandomUnderSampler(random_state=0, sampling_strategy={0: 10000, 1: 10000})
rus.fit(X_res, y_res)
X_rus, y_rus = rus.fit_resample(X_res, y_res)
print(X_rus.shape)
print(y_rus.shape)

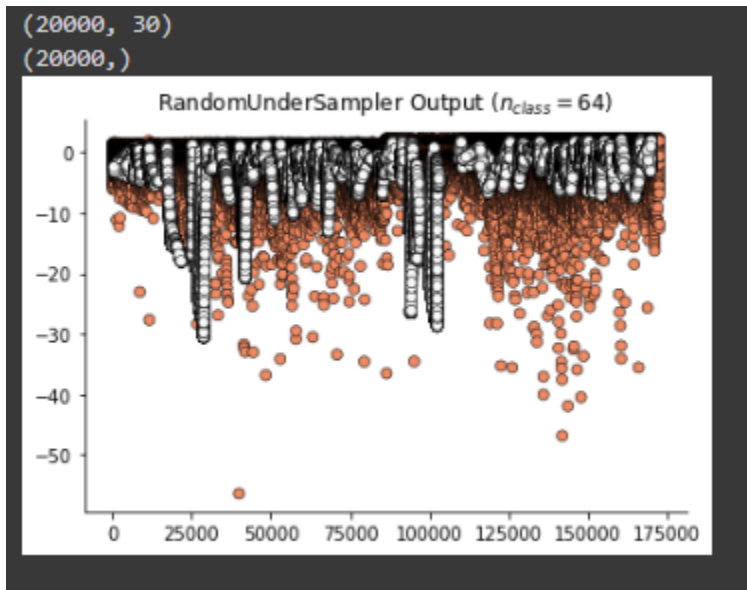
```

```

colors = ['#ef8a62' if v == 0 else '#f7f7f7' if v == 1 else '#67a9cf' for v in y_res]
plt.scatter(X_res[:, 0], X_res[:, 1], c=colors, linewidth=0.5, edgecolor='black')

```

```
sns.despine()
plt.title("RandomUnderSampler Output ($n_{class}=64)$")
pass
```

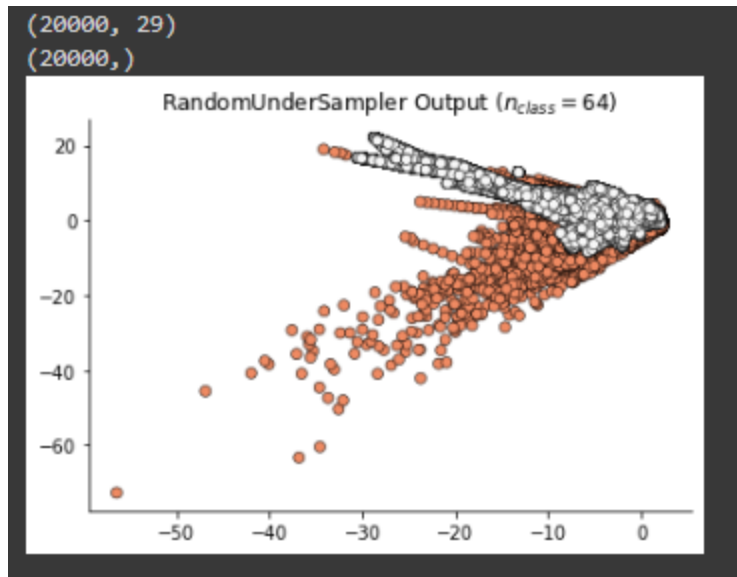


# Random Under sampling of SMOTE data

```
from imblearn.under_sampling import RandomUnderSampler
```

```
rus = RandomUnderSampler(random_state=0, sampling_strategy={0: 10000, 1: 10000})
rus.fit(x_smote, y_smote)
x_rus1, y_rus1 = rus.fit_resample(x_smote, y_smote)
print(x_rus1.shape)
print(y_rus1.shape)
```

```
colors = ['#ef8a62' if v == 0 else '#f7f7f7' if v == 1 else '#67a9cf' for v in y_smote]
plt.scatter(x_smote[:, 0], x_smote[:, 1], c=colors, linewidth=0.5, edgecolor='black')
sns.despine()
plt.title("RandomUnderSampler Output ($n_{class}=64)$")
pass
```



**Comparing the under sampled data for both the algorithms we can see that the Adasyn undersampled data has more balance compared to the SMOTE undersampled data.**

# Isolation Forest for anomaly detection

```
import numpy as np
from sklearn.ensemble import IsolationForest


random_state = np.random.RandomState(42)
model=IsolationForest(n_estimators=100,max_samples='auto',contamination=float(0.179),random_state=random_state)

model.fit(X_rus)
pred1 = model.predict(X_rus)

pred_scores1 = -1*model.score_samples(X_rus)

print(model.get_params())

plt.scatter(X_rus[:, 0], X_rus[:, 1], c=pred_scores1, cmap='RdBu')
plt.colorbar(label='Simplified Anomaly Score')
plt.show()
```



**The above plot shows lesser anomalies for the sampled data compared to the original data. Thus improving the distribution and balance of the data points.**

```
class1_count_0, class1_count_1 = y_rus.value_counts()
```

```
# Separate class
```

```
class1_0 = y_rus[y_rus == 0]
```

```
class1_1 = y_rus[y_rus == 1]# print the shape of the class
```

```
print('class 0:', class1_0.shape)
```

```
print('class 1:', class1_1.shape)
```

```
class 0: (10000,)
class 1: (10000,)
```

```
# Splitting into train and test for Adasyn data
```

```
from sklearn.model_selection import train_test_split
```

```
x_train1, x_test1, y_train1, y_test1= train_test_split(X_rus, y_rus, test_size= 0.25,
random_state=0)
```

```
# Splitting into train and test for Smote data
```

```
from sklearn.model_selection import train_test_split
```

```
x_train2, x_test2, y_train2, y_test2= train_test_split(x_rus1, y_rus1, test_size= 0.25,
random_state=0)
```

```
# Random Forest
```

```
from sklearn.ensemble import RandomForestClassifier
```

```
classifier= RandomForestClassifier(n_estimators= 10, criterion="entropy")
```

```

classifier1= RandomForestClassifier(n_estimators= 10, criterion="entropy")
classifier1.fit(x_train1, y_train1)
classifier1.fit(x_train2, y_train2)

```

```

y_pred= classifier1.predict(x_test1)
y_pred_smote=classifier1.predict(x_test2)

```

```

#Creating the Confusion matrix
from sklearn.metrics import confusion_matrix
cm= confusion_matrix(y_test1, y_pred)
print(cm)

```

```

[[2473   4]
 [  31 2492]]

```

```

from sklearn.metrics import precision_score, recall_score, accuracy_score
print(accuracy_score(y_test1, y_pred))
print(precision_score(y_test1, y_pred))
print(recall_score(y_test1, y_pred))

```

```

0.993
0.9983974358974359
0.9877130400317082

```

**The above snippet shows accuracy for Adasyn data**

```

from sklearn.metrics import precision_score, recall_score, accuracy_score
print(accuracy_score(y_test2, y_pred_smote))
print(precision_score(y_test2, y_pred_smote))
print(recall_score(y_test2, y_pred_smote))

```

```

0.9912
0.9963956748097718
0.9861276258422513

```

**The above shows accuracy for smote data**

The above two plots show adasyn has a higher accuracy compared to smote. Hence adasyn performs better compared to smote. We will be using Adasyn for all algorithms henceforth.

# KNN

```

#Fitting K-NN classifier to the training set
from sklearn.neighbors import KNeighborsClassifier
classifier= KNeighborsClassifier(n_neighbors=24, metric='minkowski', p=2 )
classifier.fit(x_train1, y_train1)

```

```
y_pred1= classifier.predict(x_test1)
```

```
from sklearn.metrics import precision_score, recall_score, accuracy_score
print(accuracy_score(y_test1, y_pred1))
print(precision_score(y_test1, y_pred1))
print(recall_score(y_test1, y_pred1))
```

```
0.872
0.9906201146430432
0.7534680935394372
```

```
# Naive Bayes
```

```
from sklearn.naive_bayes import GaussianNB
clf = GaussianNB()
clf.fit(x_train1, y_train1)
```

```
# making predictions on the testing set
y_pred2 = clf.predict(x_test1)
```

```
from sklearn.metrics import precision_score, recall_score, accuracy_score
print(accuracy_score(y_test1, y_pred2))
print(precision_score(y_test1, y_pred2))
print(recall_score(y_test1, y_pred2))
```

```
# SVM
```

```
from sklearn.svm import SVC
svclassifier = SVC(kernel='linear')
clf=svclassifier.fit(x_train1, y_train1)
```

```
y_pred3 = svclassifier.predict(x_test1)
```

```
from sklearn.metrics import precision_score, recall_score, accuracy_score
print(accuracy_score(y_test1, y_pred3))
print(precision_score(y_test1, y_pred3))
print(recall_score(y_test1, y_pred3))
```

```
0.9452
0.9870073624945864
0.9032897344431232
```

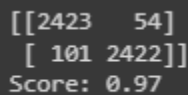
```
# Logistic regression
```

```
from sklearn.metrics import confusion_matrix
from sklearn.metrics import accuracy_score
from sklearn.metrics import classification_report
from sklearn.linear_model import LogisticRegression
```

```
classifier = LogisticRegression(random_state=0)
classifier.fit(x_train1, y_train1)
y_pred4 = classifier.predict(x_test1)
confusion_matrix = confusion_matrix(y_test1, y_pred4)
```

```
print(confusion_matrix)
```

```
print('Score: {:.2f}'.format(classifier.score(x_test1, y_test1)))
```



```
[[2423  54]
 [ 101 2422]]
Score: 0.97
```

```
# AUC ROC plots with threshold 0.5
```

```
from sklearn import metrics
import numpy as np
import matplotlib.pyplot as plt
fpr, tpr, _ = metrics.roc_curve(y_test1, y_pred)
auc = round(metrics.roc_auc_score(y_test1, y_pred), 4)
plt.plot(fpr, tpr, label="Random Forest, AUC="+str(auc))
```

```
fpr, tpr, _ = metrics.roc_curve(y_test1, y_pred1)
auc = round(metrics.roc_auc_score(y_test1, y_pred1), 4)
plt.plot(fpr, tpr, label="KNN, AUC="+str(auc))
```

```
fpr, tpr, _ = metrics.roc_curve(y_test1, y_pred2)
auc = round(metrics.roc_auc_score(y_test1, y_pred2), 4)
plt.plot(fpr, tpr, label="Naive Bayes, AUC="+str(auc))
```

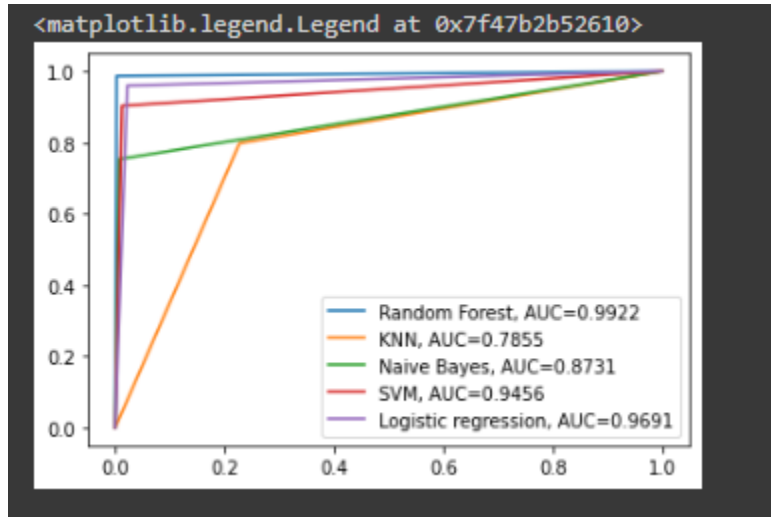
```
fpr, tpr, _ = metrics.roc_curve(y_test1, y_pred3)
auc = round(metrics.roc_auc_score(y_test1, y_pred3), 4)
plt.plot(fpr, tpr, label="SVM, AUC="+str(auc))
```

```
fpr, tpr, _ = metrics.roc_curve(y_test1, y_pred4)
```

```

auc = round(metrics.roc_auc_score(y_test1, y_pred4), 4)
plt.plot(fpr, tpr, label="Logistic regression, AUC="+str(auc))
#add legend
plt.legend()

```



**Inference:**

**1. The accuracy for different models are as follows**

**Random Forest- 99.3%**

**KNN- 78.56%**

**Naive Bayes - 87.2%**

**SVM - 94.52%**

**Logistic regression- 97%**

**2. Thus the highest accuracy occurs for random forest algorithm.**

**3. Looking at the AUC ROC plot for each graphs we can see that the data shows best results for Random Forest followed by Logistic regression and SVM**

**4. The above results for Adasyn and Smote up sampling data shows Adasyn works better for this dataset.**

**5. The imbalanced data is up sampled to balance and down sampled to reduce the size and computation power.**

**The down sampling process reduces the accuracy by a very marginal value but provides higher computer power, hence it is preferred when the data size is huge.**